

Introduction to the PETSc library
Victor Eijkhout
eijkhout@tacc.utexas.edu
2020/02/06

Introduction

To set the stage

Developing parallel, nontrivial PDE solvers that deliver high performance is still difficult and requires months (or even years) of concentrated effort. PETSc is a toolkit that can ease these difficulties and reduce the development time, but it is not black-box PDE solver, nor a silver bullet.

Barry Smith

More specifically...

Portable Extendable Toolkit for Scientific Computations

- Scientific Computations: parallel linear algebra, in particular linear and nonlinear solvers
- Toolkit: Contains high level solvers, but also the low level tools to roll your own.
- Portable: Available on many platforms, basically anything that has MPI

Why use it? It's big, powerful, well supported.

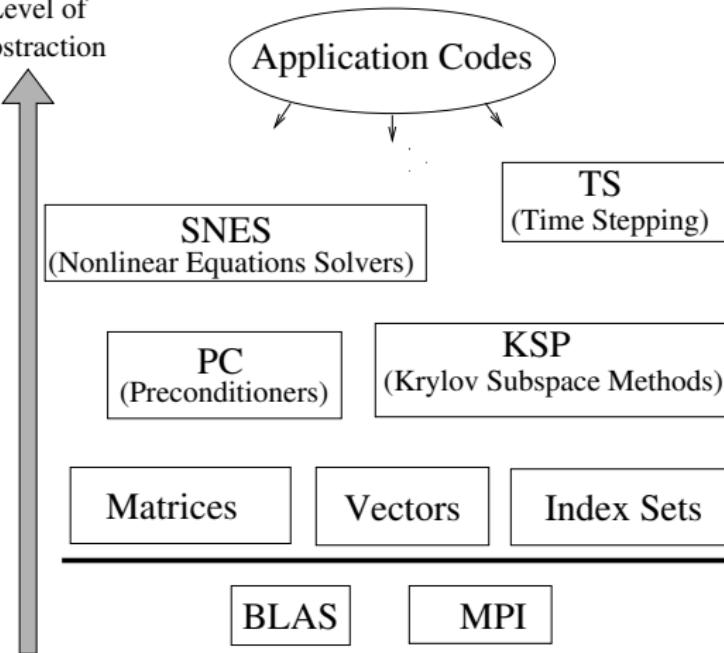
What does PETSc target?

- Serial and Parallel
- Linear and nonlinear
- Finite difference and finite element
- Structured and unstructured

What is in PETSc?

- Linear system solvers (sparse/dense, iterative/direct)
- Nonlinear system solvers
- Optimization: TAO (used to be separate library)
- Tools for distributed matrices
- Support for profiling, debugging, graphical output

Level of
Abstraction



Parallel Numerical Components of PETSc

| Nonlinear Solvers | |
|----------------------|--------------|
| Newton-based Methods | Other |
| Line Search | Trust Region |

| Time Steppers | | | |
|---------------|----------------|----------------------|-------|
| Euler | Backward Euler | Pseudo-Time Stepping | Other |
| | | | |

| Krylov Subspace Methods | | | | | | | |
|-------------------------|----|-----|------------|-------|------------|-----------|-------|
| GMRES | CG | CGS | Bi-CG-Stab | TFQMR | Richardson | Chebychev | Other |

| Preconditioners | | | | | | |
|------------------|--------------|--------|-----|-----|-------------------------|-------|
| Additive Schwarz | Block Jacobi | Jacobi | ILU | ICC | LU (sequential only) | Other |

| Matrices | | | | |
|-----------------------------|------------------------------------|------------------------|-------|-------|
| Compressed Sparse Row (AIJ) | Block Compressed Sparse Row (BAIJ) | Block Diagonal (BDiag) | Dense | Other |

| Vectors | Index Sets | | | |
|---------|------------|---------------|--------|-------|
| | Indices | Block Indices | Stride | Other |

Documentation and help

- Web page: <http://www.mcs.anl.gov/petsc/>
- PDF manual: <http://www.mcs.anl.gov/petsc/petsc-current/docs/manual.pdf>
- Follow-up to this tutorial: eijkhout@tacc.utexas.edu
- PETSc on your local cluster: ask your local support
- General questions about PETSc: petsc-maint@mcs.anl.gov
- Example codes, found online, and in
 \$PETSC_DIR/src/mat/examples et cetera
- Sometimes consult include files, for instance
 \$PETSC_DIR/include/petscmat.h

External packages

PETSc does not do everything, but it interfaces to other software:

- Dense linear algebra: Scalapack, Plapack, Elemental
- Grid partitioning software: ParMetis, Jostle, Chaco, Party
- ODE solvers: PVODE
- Eigenvalue solvers (including SVD): SLEPc

PETSc and parallelism

PETSc is layered on top of MPI

- MPI has basic tools: send elementary datatypes between processors
- PETSc has intermediate tools:
insert matrix element in arbitrary location,
do parallel matrix-vector product
- ⇒ you do not need to know much MPI when you use PETSc

PETSc and parallelism

- All objects in Petsc are defined on a communicator; can only interact if on the same communicator
- Parallelism through MPI
- Transparent: same code works sequential and parallel.
- No OpenMP used in the library:
user can use shared memory programming.
- Likewise, threading is kept outside of PETSc code: not thread-safe.
- Limited GPU! (GPU!) support; know what you're doing!
TACC note. Only available on the Frontera RTX nodes (single precision).

Object oriented design

Petsc uses objects: vector, matrix, linear solver, nonlinear solver

Overloading:

```
MATMult(A,x,y); // y <- A x
```

same for sequential, parallel, dense, sparse, FFT

Data hiding

To support this uniform interface, the implementation is hidden:

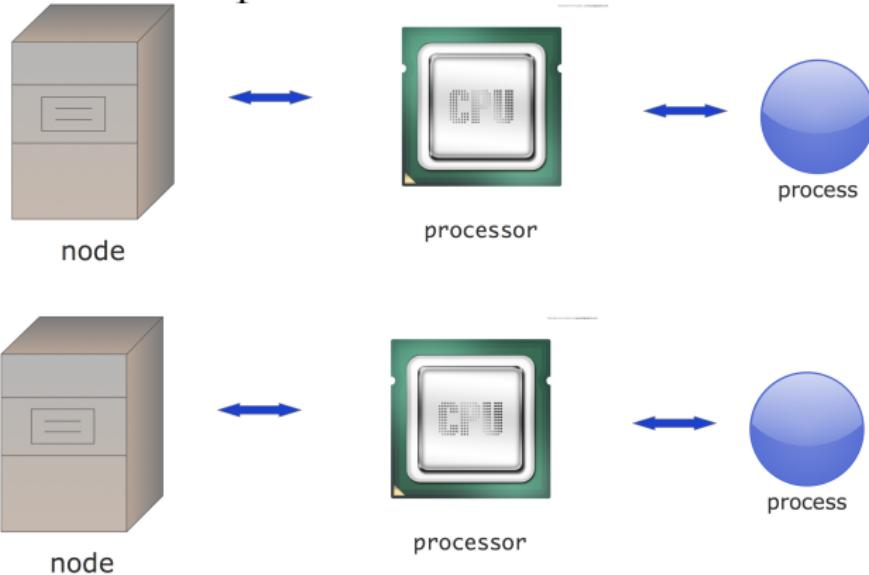
```
MatSetValue(A,i,j,v,INSERT_VALUES); // A[i,j] <- v
```

There are some direct access routines, but most of the time you don't need them.

(And don't worry about function call overhead.)

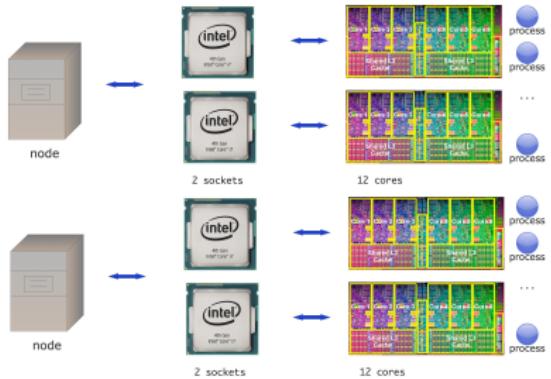
A word about SPMD parallelism

Computers when MPI was designed



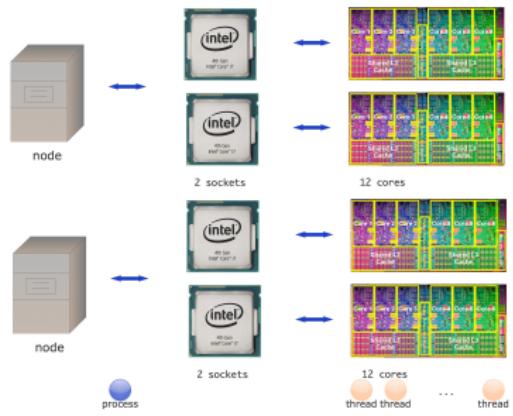
One processor and one process per node;
all communication goes through the network.

Pure MPI



A node has multiple sockets, each with multiple cores.
Pure MPI puts a process on each core: pretend shared memory
doesn't exist.

Hybrid programming



Hybrid programming puts a process per node or per socket; further parallelism comes from threading.
No use of threading in PETSc

Multi-mode programming in PETSc

PETSc is largely aimed at MPI programming; however

- You can of course use OpenMP in between PETSc calls;
- there is support for GPUs
 - TACC note. At the moment only on frontera: module load petsc/3.15-rtx.
- OpenMP can be used in external packages.

Terminology

‘Processor’ is ambiguous: is that a chip or one independent instruction processing unit?

- Socket: the processor chip
- Processor: we don’t use that word
- Core: one instruction-stream processing unit
- Process: preferred terminology in talking about MPI.

SPMD

The basic model of MPI is
‘Single Program Multiple Data’:
each process is an instance of the same program.

Symmetry: There is no ‘master process’, all processes are equal,
start and end at the same time.

Communication calls do not see the cluster structure:
data sending/receiving is the same for all neighbours.

Compiling and running

MPI compilers are usually called mpicc, mpif90, mpicxx.

These are not separate compilers, but scripts around the regular C/Fortran compiler. You can use all the usual flags.

At TACC:

ibrun yourprog

the number of processes is determined by SLURM.

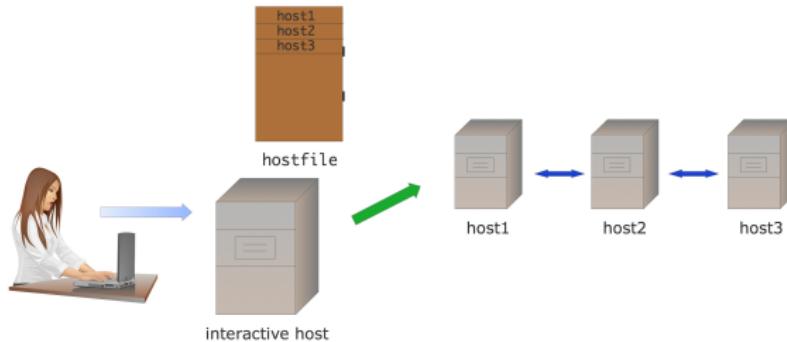
Do I need a supercomputer?

- With mpiexec and such, you start a bunch of processes that execute your PETSc program.
- Does that mean that you need a cluster or a big multicore?
- No! You can start a large number of processes, even on your laptop. The OS will use ‘time slicing’.
- Of course it will not be very efficient...

Cluster setup

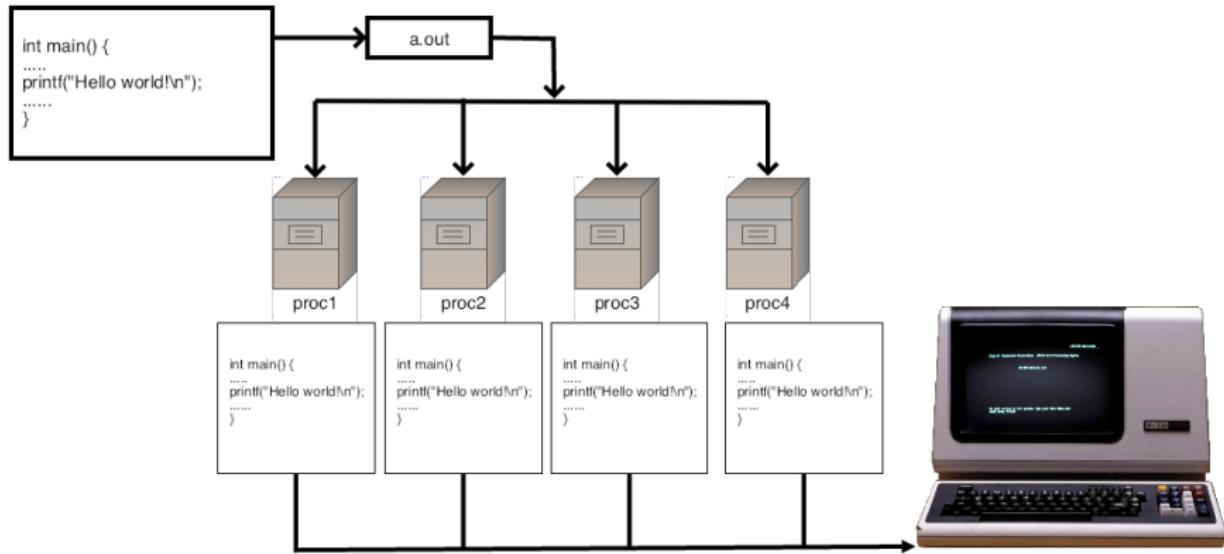
Typical cluster:

- Login nodes, where you ssh into; usually shared with 100 (or so) other people. You don't run your parallel program there!
- Compute nodes: where your job is run. They are often exclusive to you: no other users getting in the way of your program.



Hostfile: the description of where your job runs. Usually generated by a job scheduler.

In a picture



Process identification

Every process has a number (with respect to a communicator)

```
int MPI_Comm_rank( MPI_Comm comm, int *procno )
int MPI_Comm_size( MPI_Comm comm, int *nprocs )
```

For now, the communicator will be MPI_COMM_WORLD.

Note: mapping of ranks to actual processes and cores is not predictable!

Getting started

Include files

C:

```
|| #include "petsc.h"
|| int main(int argc,char **argv)
```

Fortran:

```
|| program basic
|| #include <petsc/finclude/petsc.h>
|| use petsc
|| implicit none
```

Python:

```
|| from petsc4py import PETSc
```

Variable declarations, C

```
KSP      solver;
Mat      A;
Vec      x,y;
PetscInt n = 20;
PetscScalar v;
PetscReal nrm;
```

Note Scalar vs Real

Variable declarations, F

```
KSP      :: solver
Mat      :: A
Vec      :: x,y
PetscInt :: j(3)
PetscScalar :: mv
PetscReal :: nrm
```

Much like in C; uses cpp

Library setup, C

```
// init.c
ierr = PetscInitialize(&argc,&argv,(char*)0,help); CHKERRQ(ierr);
int flag;
MPI_Initialized(&flag);
if (flag)
    printf("MPI was initialized by PETSc\n");
else
    printf("MPI not yet initialized\n");
```

Can replace MPI_Init

General: Every routine has an error return. Catch that value!

Library setup, F

```
call PetscInitialize(PETSC_NULL_CHARACTER,ierr)
      CHKERRQ(ierr)
// all the petsc work
call PetscFinalize(ierr); CHKERRQ(ierr)
```

Error code is now final parameter. This holds for every PETSc routine.

- CHKERRA in main program
- CKKERRQ in subprograms

A word about datatypes

PETSc programs can not mix single and double precision, nor real/complex:

`PetscScalar` is single/double/complex depending on the installation.

`PetscReal` is always real, even in complex installations.

Similarly, `PetscInt` is 32/64 bit depending.

Other scalar data types: `PetscBool`, `PetscErrorCode`

TACC note.

```
module spider petsc  
module avail petsc
```

```
module load petsc/3.12-i64 # et cetera
```

Debug and production

While you are developing your code:

```
module load petsc/3.15-debug
```

This does bounds tests and other time-wasting error checking.

Production:

```
module load petsc/3.15
```

This will just bomb if your program is not correct.

Every petsc configuration is available as debug and non-debug.

Exercise 1 (hello)

Look up the function PetscPrintf and print a message
‘This program runs on 27 processors’
from process zero.

- Start with the template code hello.c/hello.F
- (or see slide ??)
- Compile with make hello
- Part two: use PetscSynchronizedPrintf

PetscPrintf

About routine prototypes: C/C++

Prototype:

```
|| PetscErrorCode VecCreate(MPI_Comm comm,Vec *v);
```

Use:

```
|| PetscErrorCode ierr;
|| MPI_Comm comm = MPI_COMM_WORLD;
|| Vec v;
|| ierr = VecCreate( comm,&vec ); CHKERRQ(ierr).
```

(always good idea to catch that error code)

About routine prototypes: Fortran

Prototype

```
Subroutine VecCreate( comm,v
    ↪,ierr )
Type(MPI_Comm) :: comm
Vec           :: v
PetscErrorCode :: ierr
```

- Final parameter always error parameter. Do not forget!
- MPI types are often `Type(MPI_Comm)` and such,
- PETSc datatypes are handled through the preprocessor.

Use:

```
Type(MPI_Comm) :: comm =
    ↪MPI_COMM_WORLD
    ↪
Vec           :: v
PetscErrorCode :: ierr

call VecCreate(comm,v,ierr)
```

About routine prototypes: Python

Prototype:

```
# object method  
PETSc.Mat.setSizes(self, size, bsize=None)  
# class method  
PETSc.Sys.Print(type cls, *args, **kwargs)
```

Use:

```
from petsc4py import PETSc  
PETSc.Sys.Print("detecting n option")  
A = PETSc.Mat().create(comm=comm)  
A.setSizes( ( (None,matrix_size), (None,matrix_size) ) )
```

Note to self

```
|| PetscInitialize(&argc,&args,0,"Usage: prog -o1 v1 -o2 v2\n");
```

run as

```
./program -help
```

This displays the usage note, plus all available petsc options.

Not available in Fortran

Routine start/end, C

Debugging support:

```
|| PetscFunctionBegin;  
|| // all statements  
|| PetscFunctionReturn(0);
```

leads to informative tracebacks.

(Only in C, not in Fortran)

Example: function with error

```
// backtrace.c
PetscErrorCode this_function_bombs() {
    PetscFunctionBegin;
    SETERRQ(PETSC_COMM_SELF,1,"We cannot go on like this");
    PetscFunctionReturn(0);
}
```

Example: error traceback

```
[0]PETSC ERROR: We cannot go on like this
[0]PETSC ERROR: See https://www.mcs.anl.gov/petsc/documentation/faq.html for more information.
[0]PETSC ERROR: Petsc Release Version 3.12.2, Nov, 22, 2019
[0]PETSC ERROR: backtrace on a [computer name]
[0]PETSC ERROR: Configure options [all options]
[0]PETSC ERROR: #1 this_function_bombs() line 20 in backtrace.c
[0]PETSC ERROR: #2 main() line 30 in backtrace.c
```

Exercise 2 (root)

Start with root.c. Write a function that computes a square root, or displays an error on negative input: Look up the definition of SETERRQ1.

```
x = 1.5; ierr = square_root(x,&rootx); CHKERRQ(ierr);
PetscPrintf(PETSC_COMM_WORLD,"Root of %f is %f\n",x,rootx
           );
x = -2.6; ierr = square_root(x,&rootx); CHKERRQ(ierr);
PetscPrintf(PETSC_COMM_WORLD,"Root of %f is %f\n",x,rootx
           );
```

This should give as output:

```
Root of 1.500000 is 1.224745
[0]PETSC ERROR: ----- Error Message -----
[0]PETSC ERROR: Cannot compute the root of -2.600000
[...]
[0]PETSC ERROR: #1 square_root() line 23 in root.c
[0]PETSC ERROR: #2 main() line 39 in root.c
```

Fortran: you need to set the ierr value yourself.

Program parameters, C

(I'm leaving out the `CHKERRQ(ierr)` in the examples,
but do use this in actual code)

```
ierr = PetscOptionsGetInt  
      (PETSC_NULL,PETSC_NULL,"-n",&n,&flag); CHKERRQ(ierr);  
ierr = PetscPrintf  
      (comm,"Input parameter: %d\n",n); CHKERRQ(ierr);
```

Read commandline argument, print out from processor zero
flag can be PETSC_NULL

Program parameters, F

```
character*80      msg
call PetscOptionsGetInt(
    PETSC_NULL_OPTIONS, PETSC_NULL_CHARACTER, &
    "-n",n,PETSC_NULL_BOOL,ierr)
write(msg,10) n
10 format("Input parameter:",i5)
call PetscPrintf(PETSC_COMM_WORLD,msg,ierr)
```

PETSC_NULL_BOOL will be introduced after 3.10.2 Less elegant than PetscPrintf in C

Note the PETSC_NULL_XXX: Fortran has strict type checking.

Program parameters, Python

```
|| nlocal = PETSc.Options().getInt("n",10)
```

Vec datatype: vectors

Create calls

Everything in PETSc is an object, with create and destroy calls:

```
|| VecCreate(MPI_Comm comm,Vec *v);  
|| VecDestroy(Vec *v);  
  
|| /* C */  
|| Vec V;  
|| VecCreate(MPI_COMM_WORLD,&V);  
|| VecDestroy(&V);
```

```
|| ! Fortran  
|| Vec V  
|| call VecCreate(MPI_COMM_WORLD,V,e)  
|| call VecDestroy(V,e)
```

Note: in Fortran there are no ‘star’ arguments

Python

```
|| x = PETSc.Vec().create(comm=comm)
|| x.setType(PETSc.Vec.Type.MPI)
```

More about vectors

A vector is a vector of `PetscScalars`: there are no vectors of integers (see the `IS` datatype later)

The vector object is not completely created in one call:

```
|| VecSetType(V, VECMPI) // or VECSEQ  
|| VecSetSizes(Vec v, int m, int M);
```

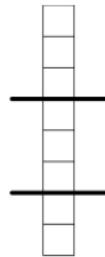
Other ways of creating: make more vectors like this one:

```
|| VecDuplicate(Vec v, Vec *w);
```

Parallel layout up to PETSc

```
|| VecSetSizes(Vec v, int m, int M);
```

Local size can be specified as `PETSC_DECIDE`.



```
VecSetSizes(V,PETSC_DECIDE,8)
```

```
VecSetSizes(V,PETSC_DECIDE,8)
```

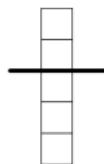
```
VecSetSizes(V,PETSC_DECIDE,8)
```

Parallel layout specified

Local or global size in

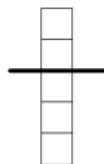
```
|| VecSetSizes(Vec v, int m, int M);
```

Global size can be specified as PETSC_DECIDE.



VecSetSizes(V,2,5)

VecSetSizes(V,3,5)



VecSetSizes(V,2,PETSC_DECIDE)

VecSetSizes(V,3,PETSC_DECIDE)

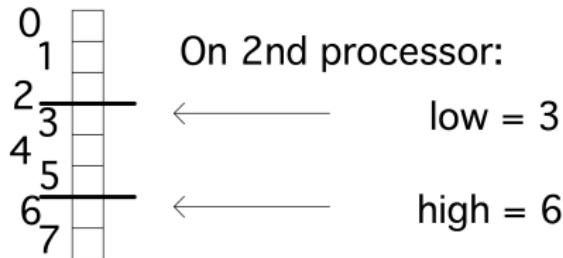
Python

```
|| x.setSizes([n,PETSc.DECIDE])  
|| x.setFromOptions()
```

Query parallel layout

Query vector layout:

```
|| VecGetSize(Vec,PetscInt *globalsize)
|| VecGetLocalSize(Vec,PetscInt *localsize)
|| VecGetOwnershipRange(Vec x,PetscInt *low,PetscInt *high)
|| VecGetOwnershipRange(x,low,high,ierr) ! F
```



Layout, regardless object

Query general layout:

```
|| PetscSplitOwnership(MPI_Comm comm,PetscInt *n,PetscInt *N);  
|| PetscSplitOwnership(comm,n,N,ierr) ! F
```

(get local/global given the other)

Setting values

Set vector to constant value:

```
|| VecSet(Vec x,PetscScalar value);
```

Set individual elements (global indexing!):

```
|| VecSetValue(Vec x,int row,PetscScalar value, InsertMode mode);
```

```
|| i = 1; v = 3.14;
```

```
|| VecSetValue(x,i,v,INSERT_VALUES);
```

```
|| call VecSetValue(x,i,v,INSERT_VALUES,ierr,e)
```

Setting values by block

Set individual elements (global indexing!):

```
|| VecSetValues(Vec x,int n,int *rows,PetscScalar *values,  
|| InsertMode mode); // INSERT_VALUES or ADD_VALUES  
  
|| ii[0] = 1; ii[1] = 2; vv[0] = 2.7; vv[1] = 3.1;  
|| VecSetValues(x,2,ii,vv,INSERT_VALUES);  
  
|| ii(1) = 1; ii(2) = 2; vv(1) = 2.7; vv(2) = 3.1  
|| call VecSetValues(x,2,ii,vv,INSERT_VALUES,ierr,e)
```

Setting values

No restrictions on parallelism;
after setting, move values to appropriate processor:

```
|| VecAssemblyBegin(Vec x);  
|| VecAssemblyEnd(Vec x);
```

Basic operations

```
VecAXPY(Vec y,PetscScalar a,Vec x); /* y <- y + a x */
VecAYPX(Vec y,PetscScalar a,Vec x); /* y <- a y + x */
VecScale(Vec x, PetscScalar a);
VecDot(Vec x, Vec y, PetscScalar *r); /* several variants */
VecMDot(Vec x,int n,Vec y[],PetscScalar *r);
VecNorm(Vec x,NormType type, PetscReal *r);
VecSum(Vec x, PetscScalar *r);
VecCopy(Vec x, Vec y);
VecSwap(Vec x, Vec y);
VecPointwiseMult(Vec w,Vec x,Vec y);
VecPointwiseDivide(Vec w,Vec x,Vec y);
VecMAXPY(Vec y,int n, PetscScalar *a, Vec x[]);
VecMax(Vec x, int *idx, double *r);
VecMin(Vec x, int *idx, double *r);
VecAbs(Vec x);
VecReciprocal(Vec x);
VecShift(Vec x,PetscScalar s);
```

Exercise 3 (vec)

Create a vector where the values are a single sine wave. using
`VecGetSize`, `VecGetLocalSize`, `VecGetOwnershipRange`. Quick visual
inspection:

```
iBrun vec -n 12 -vec_view
```

Exercise 4 (vec)

Use the routines `VecDot`, `VecScale` and `VecNorm` to compute the inner product of vectors x, y , scale the vector x , and check its norm:

$$\begin{aligned} p &\leftarrow x^t y \\ x &\leftarrow x/p \\ n &\leftarrow \|x\|_2 \end{aligned}$$

Split dot products and norms

MPI is capable (in principle) of ‘overlapping computation and communication’.

- Start inner product / norm with `VecDotBegin` / `VecNormBegin`;
- Conclude inner product / norm with `VecDotEnd` / `VecNormEnd`;

Also: start/end multiple norm/dotproduct operations.

Direct access to vector values (C)

Setting values is done without user access to the stored data
Getting values is often not necessary: many operations provided.

what if you do want access to the data?

Solution 1. Create vector from user provided array:

```
|| VecCreateSeqWithArray(MPI_Comm comm,  
||   PetscInt n,const PetscScalar array[],Vec *V)  
|| VecCreateMPIWithArray(MPI_Comm comm,  
||   PetscInt n,PetscInt N,const PetscScalar array[],Vec *vv)
```

Direct access'

Solution 2. Retrive the internal array:

```
|| VecGetArray(Vec x,PetscScalar *a[])
/* do something with the array */
|| VecRestoreArray(Vec x,PetscScalar *a[])
```

Note: local only; see [VecScatter](#) for more general mechanism)

Getting values example

```
int localsize,first,i;
PetscScalar *a;
VecGetLocalSize(x,&localsize);
VecGetOwnershipRange(x,&first,PETSC_NULL);
VecGetArray(x,&a);
for (i=0; i<localsize; i++)
    printf("Vector element %d : %e\n",first+i,a[i]);
VecRestoreArray(x,&a);
```

Fortran: PETSC_NULL_INTEGER

More array juggling

- `VecPlaceArray`: replace the internal array; the original can be restored with `VecRestoreArray`
- `VecReplaceArray`: replace and free the internal array.

Array handling in F90

```
PetscScalar, pointer :: xx_v(:)  
....  
call VecGetArrayF90(x,xx_v,ierr)  
a = xx_v(3)  
call VecRestoreArrayF90(x,xx_v,ierr)
```

More separate F90 versions for ‘Get’ routines
(there are some ugly hacks for F77)

Mat Datatype: matrix

Matrix creation

The usual create/destroy calls:

```
|| MatCreate(MPI_Comm comm,Mat *A)  
|| MatDestroy(Mat *A)
```

Several more aspects to creation:

```
|| MatSetType(A,MATSEQAIJ) /* or MATMPIAIJ or MATAIJ */  
|| MatSetSizes(Mat A,int m,int n,int M,int N)  
|| MatSeqAIJSetPreallocation /* more about this later*/  
||   (Mat B,PetscInt nz,const PetscInt nnz[])
```

Local or global size can be PETSC_DECIDE (as in the vector case)

If you already have a CRS matrix

```
PetscErrorCode MatCreateSeqAIJWithArrays  
  (MPI_Comm comm,PetscInt m,PetscInt n,  
   PetscInt* i,PetscInt*j,PetscScalar *a,Mat *mat)
```

(also from triplets)

Do not use this unless you interface to a legacy code. And even then...

Matrix Preallocation

- PETSc matrix creation is very flexible:
- No preset sparsity pattern
- any processor can set any element
⇒ potential for lots of malloc calls
- tell PETSc the matrix' sparsity structure
(do construction loop twice: once counting, once making)
- Re-allocating is expensive:

```
|| MatSetOption(A,MAT_NEW_NONZERO_LOCATIONS,  
    ↪PETSC_FALSE);
```

(is default) Otherwise:

[1]PETSC ERROR: Argument out of range

[1]PETSC ERROR: New nonzero at (0,1) caused a malloc

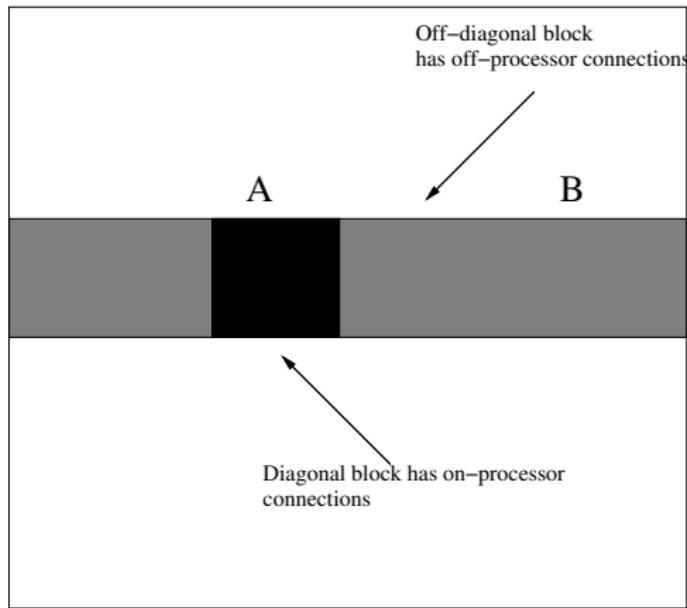
Sequential matrix structure

```
|| MatSeqAIJSetPreallocation  
||   (Mat B,PetscInt nz,const PetscInt nnz[])
```

- nz number of nonzeros per row
(or slight overestimate)
- nnz array of row lengths (or overestimate)
- considerable savings over dynamic allocation!

In Fortran use `PETSC_NULL_INTEGER` if not specifying nnz array

Parallel matrix structure



(why does it do this?)

- $y \leftarrow Ax_A + Bx_B$
- x_B needs to be communicated; Ax_A can be computed in the meantime
- Algorithm
 - Initiate asynchronous sends/receives for x_B
 - compute Ax_A
 - make sure x_B is in
 - compute Bx_B
- so by splitting matrix storage into A,B part, code for the sequential case can be reused.
- This is one of the few places where PETSc's design is visible to the user.

Parallel matrix structure description

- m,n local size; M,N global. Note: If the matrix is square, specify m,n equal, even though distribution by block rows
- d_nz: number of nonzeros per row in diagonal part
- o_nz: number of nonzeros per row in off-diagonal part
- d_nnz: array of numbers of nonzeros per row in diagonal part
- o_nnz: array of numbers of nonzeros per row in off-diagonal part

|| MatMPIAIJSetPreallocation

(Mat B,
 PetscInt d_nz,const PetscInt d_nnz[],
 PetscInt o_nz,const PetscInt o_nnz[])

In Fortran use `PETSC_NULL_INTEGER` if not specifying arrays

Matrix creation all in one

```
MatCreateSeqAIJ(MPI_Comm comm,PetscInt m,PetscInt n,  
    PetscInt nz,const PetscInt nnz[],Mat *A)  
MatCreateMPIAIJ(MPI_Comm comm,  
    PetscInt m,PetscInt n,PetscInt M,PetscInt N,  
    PetscInt d_nz,const PetscInt d_nnz[],  
    PetscInt o_nz,const PetscInt o_nnz[],  
    Mat *A)
```

Querying parallel structure

Matrix partitioned by block rows:

```
|| MatGetSize(Mat mat,PetscInt *M,PetscInt* N);  
|| MatGetLocalSize(Mat mat,PetscInt *m,PetscInt* n);  
|| MatGetOwnershipRange(Mat A,int *first row,int *last row);
```

In query functions, unneeded components can be specified as **PETSC_NULL**.

Fortran: **PETSC_NULL_INTEGER**

Setting values

Set one value:

```
|| MatSetValue(Mat A,  
    PetscInt i,PetscInt j,PetscScalar va,InsertMode mode)
```

where insert mode is `INSERT_VALUES`, `ADD_VALUES`

Set block of values:

```
|| MatSetValues(Mat A,int m,const int idxm[],  
    int n,const int idxn[],const PetscScalar values[],  
    InsertMode mode)
```

(v is row-oriented)

Special case of the general case:

```
|| MatSetValues(A,1,&i,1,&j,&v,INSERT_VALUES); // C  
|| MatSetValues(A,1,i,1,j,v,INSERT_VALUES,e); ! F
```

Assembling the matrix

Setting elements is independent of parallelism; move elements to proper processor:

```
|| MatAssemblyBegin(Mat A,MAT_FINAL_ASSEMBLY);  
|| MatAssemblyEnd(Mat A,MAT_FINAL_ASSEMBLY);
```

Cannot mix inserting/adding values: need to do assembly in between with **MAT_FLUSH_ASSEMBLY**

Exercise 5 (matvec)

Pretend that you do not know how the matrix is created. Use `MatGetOwnershipRange` or `MatGetLocalSize` to create a vector with the same distribution, and then compute $y \leftarrow Ax$.

(Part of the code has been disabled with `#if 0`. We will get to that next.)

Getting values (C)

- Values are often not needed: many matrix operations supported
- Matrix elements can only be obtained locally.

```
|| PetscErrorCode MatGetRow(Mat mat,  
|| PetscInt row,PetscInt *ncols,const PetscInt *cols[],  
|| const PetscScalar *vals[])  
|| PetscErrorCode MatRestoreRow(/* same parameters */)
```

Note: for inspection only; possibly expensive.

Getting values (F)

- || `MatGetRow(A,row,ncols,cols,vals,ierr)`
- || `MatRestoreRow(A,row,ncols,cols,vals,ierr)`

where `cols(maxcols)`, `vals(maxcols)` are long enough arrays
(allocated by the user)

Exercise 6 (matvec)

Advanced exercise: create a sequential (uni-processor) vector.
Question: how does the code achieve this? Give it the data of the distributed vector. Use that to compute the vector norm on each process separately.

(Start by removing the #if 0 and #endif.)

Other matrix types

MATBAIJ : blocked matrices (dof per node)

(see PETSC_DIR/include/petscmat.h)

Dense:

```
|| MatCreateSeqDense(PETSC_COMM_SELF,int m,int n,  
||   PetscScalar *data,Mat *A);  
|| MatCreateDense(MPI_Comm comm,  
||   PetscInt m,PetscInt n,PetscInt M,PetscInt N,  
||   PetscScalar *data,Mat *A)  
|| fg
```

Data argument optional: **PETSC_NULL** or
PETSC_NULL_SCALAR causes allocation

Matrix operations

Main operations are matrix-vector:

- || MatMult(Mat A, Vec in, Vec out);
- || MatMultAdd
- || MatMultTranspose
- || MatMultTransposeAdd

Simple operations on matrices:

- || MatNorm
- || MatScale
- || MatDiagonalScale

Some matrix-matrix operations

```
MatMatMult(Mat,Mat,MatReuse,PetscReal,Mat*);
```

```
MatPtAP(Mat,Mat,MatReuse,PetscReal,Mat*);
```

```
MatMatMultTranspose(Mat,Mat,MatReuse,PetscReal,Mat*);
```

```
MatAXPY(Mat,PetscScalar,Mat,MatStructure);
```

Matrix viewers

```
|| MatView(A,PETSC_VIEWER_STDOUT_WORLD);
|| row 0: (0, 1)  (2, 0.333333)  (3, 0.25)  (4, 0.2)
|| row 1: (0, 0.5)  (1, 0.333333)  (2, 0.25)  (3, 0.2)
|| ....
```

(Fortran: PETSC_NULL_INTEGER)

- also invoked by -mat_view
- Sparse: only allocated positions listed
- other viewers: for instance -mat_view_draw (X terminal)

General viewers

Any PETSc object can be ‘viewed’

- Terminal output: useful for vectors and matrices but also for solver objects.
- Binary output: great for vectors and matrices.
- Viewing can go both ways: load a matrix from file or URL into an object.
- Viewing through a socket, to Matlab or Mathematica, HDF5, VTK.

```
PetscViewer fd;  
PetscViewerCreate( comm, &fd );  
PetscViewerSetType( fd,PETSCVIEWERVTK );  
MatView( A,fd );  
PetscViewerDestroy(fd);
```

Shell matrices

What if the matrix is a user-supplied operator, and not stored?

```
MatSetType(A,MATSHELL); /* or */  
MatCreateShell(MPI Comm comm,  
    int m,int n,int M,int N,void *ctx,Mat *mat);  
  
PetscErrorCode UserMult(Mat mat,Vec x,Vec y);  
  
MatShellSetOperation(Mat mat,MatOperation MATOP_MULT,  
    (void(*)(void)) PetscErrorCode (*UserMult)(Mat,Vec,Vec));
```

Inside iterative solvers, PETSc calls `MatMult(A,x,y)`:
no difference between stored matrices and shell matrices

Shell matrix context

Shell matrices need custom data

```
|| MatShellSetContext(Mat mat,void *ctx);  
|| MatShellGetContext(Mat mat,void **ctx);
```

(This does not work in Fortran: use Common or Module or write interface block)

User program sets context, matmult routine accesses it

Shell matrix example

```
...
MatSetType(A,MATSHELL);
MatShellSetOperation(A,MATOP_MULT,(void*)&mymatmult);
MatShellSetContext(A,(void*)&mystruct);
...

PetscErrorCode mymatmult(Mat mat,Vec in,Vec out)
{
    PetscFunctionBegin;
    MatShellGetContext(mat,(<void>**)&mystruct);
    /* compute out from in, using mystruct */
    PetscFunctionReturn(0);
}
```

Submatrices

Extract one parallel submatrix:

```
|| MatGetSubMatrix(Mat mat,  
|| IS isrow,IS iscol,PetscInt csize,MatReuse cll,  
|| Mat *newmat)
```

Extract multiple single-processor matrices:

```
|| MatGetSubMatrices(Mat mat,  
|| PetscInt n,const IS irow[],const IS icol[],MatReuse scall,  
|| Mat *submat[])
```

Collective call, but different index sets per processor

Load balancing

```
|| MatPartitioningCreate  
||   (MPI Comm comm,MatPartitioning *part);
```

Various packages for creating better partitioning: Chaco, Parmetis

KSP & PC: Iterative solvers

What are iterative solvers?

Solving a linear system $Ax = b$ with Gaussian elimination can take lots of time/memory.

Alternative: iterative solvers use successive approximations of the solution:

- Convergence not always guaranteed
- Possibly much faster / less memory
- Basic operation: $y \leftarrow Ax$ executed once per iteration
- Also needed: preconditioner $B \approx A^{-1}$

Topics

- All linear solvers in PETSc are iterative, even the direct ones
- Preconditioners
- Fargoeing control through commandline options
- Tolerances, convergence and divergence reason
- Custom monitors and convergence tests

Iterative solver basics

```
KSPCreate(comm,&solver); KSPDestroy(solver);
// set Amat and Pmat
KSPSetOperators(solver,A,B); // usually: A,A
// solve
KSPSolve(solver,rhs,sol);
```

Optional: KSPSetup(solver)

Reuse Amat or Pmat: KSPGetOperators and
PetscObjectReference.

Solver type

```
KSPSetType(solver,KSPGMRES);
```

KSP can be controlled from the commandline:

```
KSPSetFromOptions(solver);
/* right before KSPSolve or KSPSetUp */
```

then options -ksp.... are parsed.

- type: -ksp_type gmres -ksp_gmres_restart 20
- -ksp_view

Convergence

Iterative solvers can fail

- Solve call itself gives no feedback: solution may be completely wrong
- `KSPGetConvergedReason(solver,&reason)` : positive is convergence, negative divergence
`KSPConvergedReasons[reason]` is string
- `KSPGetIterationNumber(solver,&nits)` : after how many iterations did the method stop?

```
KSPSolve(solver,B,X);
KSPGetConvergedReason(solver,&reason);
if (reason<0) {
    PetscPrintf(comm,
        "Failure to converge after %d iterations; reason %s\n",
        its,KSPConvergedReasons[reason]);
} else {
    KSPGetIterationNumber(solver,&its);
    PetscPrintf(comm,
        "Convergence in %d iterations.\n",its);
}
```

Preconditioners

System $Ax = b$ is transformed:

$$M^{-1}A = M^{-1}b$$

- M is constructed once, applied in every iteration
- If $M = A$: convergence in one iteration
- Tradeoff: M expensive to construct \Rightarrow low number of iterations; construction can sometimes be amortized.
- Other tradeoff: M more expensive to apply and only modest decrease in number of iterations
- Symmetry: A, M symmetric $\not\Rightarrow M^{-1}A$ symmetric, however can be symmetrized by change of inner product
- Can be tricky to make both parallel and efficient

PC basics

- PC usually created as part of KSP: separate create and destroy calls exist, but are (almost) never needed
 - KSP solver; PC precon;
 - KSPCreate(comm,&solver);
 - KSPGetPC(solver,&precon);
 - PCSetType(precon,PCJACOBI);
- Many choices, some with options: PCJACOBI, PCILU (only sequential), PCASM, PCBjacobi, PCMG, et cetera
- Controllable through commandline options:
-pc_type ilu -pc_factor_levels 3

Preconditioner reuse

In context of nonlinear solvers, the preconditioner can sometimes be reused:

- If the jacobian doesn't change much, reuse the preconditioner completely
- If the preconditioner is recomputed, the sparsity pattern probably stays the same

`KSPSetOperators(solver,A,B)`

- B is basis for preconditioner, need not be A
- if A or B is to be reused, use `NULL`

Types of preconditioners

- Simple preconditioners: Jacobi, SOR, ILU
- Compose simple preconditioners:
 - composing in space: Block Jacobi, Schwarz
 - composing in physics: Fieldsplit
- Global parallel preconditioners: multigrid, approximate inverses

Simple preconditioners

$$A = D_A + L_A + U_A, M = \dots$$

- None: $M = I$
- Jacobi: $M = D_A$
 - very simple, better than nothing
 - Watch out for zero diagonal elements
- Gauss-Seidel: $M = D_A + L_A$
 - Non-symmetric
 - popular as multigrid smoother
- SOR: $M = \omega^{-1}D_A + L_A$
 - estimating ω often infeasible
- SSOR: $M = (I + (\omega^{-1}D_A)^{-1} + L_A)(\omega^{-1}D_A + U_A)$

Factorization preconditioners

Exact factorization: $A = LU$

Inexact factorization: $A \approx M = LU$ where L, U obtained by throwing away ‘fill-in’ during the factorization process.

Exact:

$$\forall_{i,j} : a_{ij} \leftarrow a_{ij} - a_{ik} a_{kk}^{-1} a_{kj}$$

Inexact:

$$\forall_{i,j} : \text{if } a_{ij} \neq 0 \quad a_{ij} \leftarrow a_{ij} - a_{ik} a_{kk}^{-1} a_{kj}$$

Application of the preconditioner (that is, solve $Mx = y$) approx same cost as matrix-vector product $y \leftarrow Ax$

Factorization preconditioners are sequential

ILU

PCICC: symmetric, PCILU: nonsymmetric
many options:

```
PCFactorSetLevels(PC pc,int levels);  
-pc_factor_levels <levels>
```

Prevent indefinite preconditioners:

```
PCFactorSetShiftType(PC pc,MatFactorShiftType type);
```

value MAT_SHIFT_POSITIVE_DEFINITE et cetera

Block Jacobi and Additive Schwarz, theory

- Both methods parallel
- Jacobi fully parallel
 - Schwarz local communication between neighbours
- Both require sequential local solver: composition with simple preconditioners
- Jacobi limited reduction in iterations
 - Schwarz can be optimal

Block Jacobi and Additive Schwarz, coding

```
KSP *ksps; int nlocal,firstlocal; PC pc;  
PCBJacobiGetSubKSP(pc,&nlocal,&firstlocal,&ksps);  
for (i=0; i<nlocal; i++) {  
    KSPSetType( ksps[i], KSPGMRES );  
    KSPGetPC( ksps[i], &pc );  
    PCSetType( pc, PCILU );  
}
```

Much shorter: commandline options -sub_ksp_type and
-sub_pc_type (subksp is PREONLY by default)

```
PCASMSetOverlap(PC pc,int overlap);
```

Exercise 7 (ksp)

File ksp.c / ksp.F90 contains the solution of a (possibly nonsymmetric) linear system.

Compile the code and run it. Now experiment with commandline options. Make notes on your choices and their outcomes.

- The code has two custom commandline switch:
 - -n 123 set the domain size to 123 and therefore the matrix size to 123^2 .
 - -unsymmetry 456 adds a convection-like term to the matrix, making it unsymmetric. The numerical value is the actual element size that is set in the matrix.
- What is the default solver in the code? Run with -ksp_view
- Print out the matrix for a small size with -mat_view.
- Now out different solvers for different matrix sizes and amounts of unsymmetry. See the instructions in the code!

Exercise 8 (shell)

After the main program, a routine mymatmult is declared, which is attached by MatShellSetOperation to the matrix A as the means of computing the product MatMult(A,in,out), for instance inside an iterative method.

In addition to the shell matrix A, the code also creates a traditional matrix AA. Your assignment is to make it so that mymatmult computes the product $y \leftarrow A^t Ax$.

In C, use MatShellSetContext to attach AA to A and MatShellGetContext to retrieve it again for use; in Fortran use a common block (or a module) to store AA.

The code uses a preconditioner PCNONE. What happens if you run it with option -pc_type jacobi?

Monitors and convergence tests

```
KSPSetTolerances(solver,rtol,atol,dtol,maxit);
```

Monitors can be set in code, but simple cases:

- -ksp_monitor
- -ksp_monitor_true_residual

Custom monitors and convergence tests

```
KSPMonitorSet(KSP ksp,  
PetscErrorCode (*monitor)  
    (KSP,PetscInt,PetscReal,void*),  
void *mctx,  
PetscErrorCode (*monitordestroy)(void*));  
KSPSetConvergenceTest(KSP ksp,  
PetscErrorCode (*converge)  
    (KSP,PetscInt,PetscReal,KSPConvergedReason*,void*),  
void *cctx,  
PetscErrorCode (*destroy)(void*))
```

Example of convergence tests

```
PetscErrorCode resconverge
(KSP solver,PetscInt it,PetscReal res,
 KSPConvergedReason *reason,void *ctx)
{
    MPI_Comm comm; Mat A; Vec X,R; PetscErrorCode ierr;
    PetscFunctionBegin;
    KSPGetOperators(solver,&A,PETSC_NULL,PETSC_NULL);
    PetscObjectGetComm((PetscObject)A,&comm);
    KSPBuildResidual(solver,PETSC_NULL,PETSC_NULL,&R);
    KSPBuildSolution(solver,PETSC_NULL,&X);
    /* stuff */
    if (sometest) *reason = 15;
    else *reason = KSP_CONVERGED_ITERATING;
    PetscFunctionReturn(0);
}
```

Advanced options

Many options for the (mathematically) sophisticated user
some specific to one method

KSPSetInitialGuessNonzero

KSPGMRESSetRestart

KSPSetPreconditionerSide

KSPSetNormType

Many options easier through commandline.

Null spaces

```
MatNullSpace sp;  
MatNullSpaceCreate /* constant vector */  
  (PETSC_COMM_WORLD,PETSC_TRUE,0,PETSC_NULL,&sp);  
MatNullSpaceCreate /* general vectors */  
  (PETSC_COMM_WORLD,PETSC_FALSE,5,vecs,&sp);  
KSPSetNullSpace(ksp,sp);
```

The solver will now properly remove the null space at each iteration.

Matrix-free solvers

Shell matrix requires shell preconditioner in KSPSetOperators):

```
PCSetType(pc,PCSHELL);
PCShellSetContext(PC pc,void *ctx);
PCShellGetContext(PC pc,void **ctx);
PCShellSetApply(PC pc,
    PetscErrorCode (*apply)(void*,Vec,Vec));
PCShellSetSetUp(PC pc,
    PetscErrorCode (*setup)(void*))
```

similar idea to shell matrices

Alternative: use different operator for preconditioner

Fieldsplit preconditioners

If a problem contains multiple physics, separate preconditioning can make sense

Matrix block storage: MatCreateNest

$$\begin{pmatrix} A_{00} & A_{01} & A_{02} \\ A_{10} & A_{11} & A_{12} \\ A_{20} & A_{21} & A_{22} \end{pmatrix}$$

However, it makes more sense to interleave these fields

Easy case: all fields are the same size

```
PCSetType(prec,PCFIELDSPLIT);
PCFieldSplitSetBlockSize(prec,3);
PCFieldSplitSetType(prec,PC_COMPOSITE_ADDITIVE);
```

Subpreconditioners can be specified in code, but easier with options:

```
PetscOptionsSetValue
("-fieldsplit_0_pc_type","lu");
PetscOptionsSetValue
("-fieldsplit_0_pc_factor_mat_solver_package","mumps");
```

Fields can be named instead of numbered.

Non-strided, arbitrary fields: PCFieldSplitSetIS()

Stokes equation can be detected:

-pc_fieldsplit_detect_saddle_point

Combining fields multiplicatively: solve

$$\begin{pmatrix} I & \\ A_{10}A_{00}^{-1} & I \end{pmatrix} \begin{pmatrix} A_{00} & A_{01} \\ & A_{11} \end{pmatrix}$$

If there are just two fields, they can be combined by Schur complement

$$\begin{pmatrix} I & \\ A_{10}A_{00}^{-1} & I \end{pmatrix} \begin{pmatrix} A_{00} & A_{01} \\ & A_{11} - A_{10}A_{00}^{-1}A_{01} \end{pmatrix}$$

Fieldsplit example

```
KSPGetPC(solver,&prec);
PCSetType(prec,PCFIELDSPLIT);
PCFieldSplitSetBlockSize(prec,2);
PCFieldSplitSetType(prec,PC_COMPOSITE_ADDITIVE);
PetscOptionsSetValue
    ("-fieldsplit_0_pc_type","lu");
PetscOptionsSetValue
    ("-fieldsplit_0_pc_factor_mat_solver_package","mumps");
PetscOptionsSetValue
    ("-fieldsplit_1_pc_type","lu");
PetscOptionsSetValue
    ("-fieldsplit_1_pc_factor_mat_solver_package","mumps");
```

Global preconditioners: MG

```
PCSetType(PC pc,PCMGMG);
PCMGSetsLevels(pc,int levels,MPI Comm *comms);
PCMGSetsType(PC pc,PCMGTyoe mode);
PCMGSetsCycleType(PC pc,PCMGCycleTyoe ctype);
PCMGSetsNumberSmoothUp(PC pc,int m);
PCMGSetsNumberSmoothDown(PC pc,int n);
PCMGGetsCoarseSolve(PC pc,KSP *ksp);
PCMGSetsInterpolation(PC pc,int level,Mat P); and
PCMGSetsRestriction(PC pc,int level,Mat R);
PCMGSetsResidual(PC pc,int level,PetscErrorCode
(*residual)(Mat,Vec,Vec,Vec),Mat mat);
```

Global preconditioners: Hypre

- Hypre is a package like PETSc
- selling point: fancy preconditioners
- Install with `-with-hypre=yes` `-download-hypre=yes`
- then use `-pc_type hypre` `-pc_hypre_type`
`parasails`/`boomeramg`/`euclid`/`pilut`

Direct methods

- Iterative method with direct solver as preconditioner would converge in one step
- Direct methods in PETSc implemented as special iterative method: KSPPREONLY only apply preconditioner
- All direct methods are preconditioner type PCLU:

```
myprog -pc_type lu -ksp_type preonly \
-pc_factor_mat_solver_package mumps
```

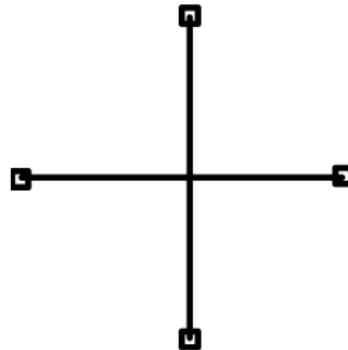
Grid manipulation

Regular grid: DMDA

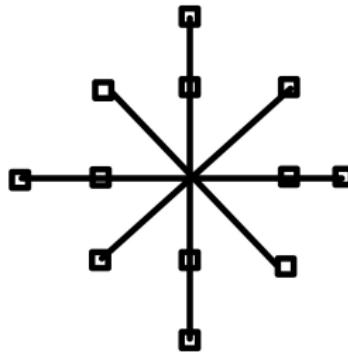
DMDAs are for storing vector field, not matrix.

Support for different stencil types:

Star stencil

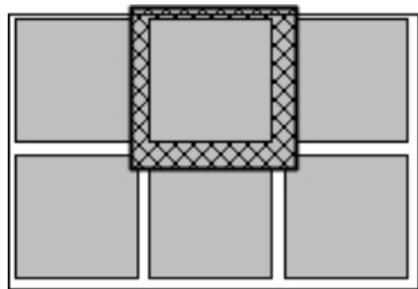


Box stencil



Ghost regions around processors

A DM_A defines a global vector, which contains the elements of the grid, and a local vector for each processor which has space for "ghost points".



DMDA construction

```
|| DMDACreate2d(comm, bndx,bndy, type, M, N, m, n,  
||   dof, s, lm[], ln[], DMDA *da)
```

bndx,bndy boundary behaviour: none/ghost/periodic

type: Specifies stencil

DMDA_STENCIL_BOX or DMDA_STENCIL_STAR

M/N: Number of grid points in x/y-direction

m/n: Number of processes in x/y-direction

dof: Degrees of freedom per node

s: The stencil width (for instance, 1 for 2D five-point stencil)

lm/n: array of local sizes (optional; Use PETSC_NULL for the default)

Associated vectors

```
DMCreateGlobalVector(DM grid,Vec *g);  
DMCreateLocalVector(DM grid,Vec *l);
```

global -> local

```
DMGlobalToLocalBegin/End  
(DMDA da,Vec g,InsertMode iora,Vec l);
```

local -> global

```
DMLocalToGlobalBegin/End  
(DMDA da,Vec l,InsertMode mode,Vec g);
```

local -> global -> local :

```
DMLocalToLocalBegin/End  
(DMDA da,Vec l1,InsertMode iora,Vec l2);
```

Associated matrix

Matrix that has knowledge of the grid:

```
|| DMSetUp(DM grid);
|| DMCreatematrix(DM grid,Mat *J)
```

Set matrix values based on stencil:

```
|| MatSetValuesStencil(Mat mat,
|| PetscInt m,const MatStencil idxm[],
|| PetscInt n,const MatStencil idxn[],
|| const PetscScalar v[],InsertMode addv)
```

(ordering of row/col variables too complicated for
MatSetValues)

Grid info

```
typedef struct {
    PetscInt      dim,dof,sw;
    PetscInt      mx,my,mz; /* grid points in x,y,z */
    PetscInt      xs,ys,zs; /* starting point, excluding ghosts */
    PetscInt      xm,ym,zm; /* grid points, excluding ghosts */
    PetscInt      gxs,gys,gzs; /* starting point, including ghosts */
    PetscInt      gxm,gym,gzm; /* grid points, including ghosts */
    DMBoundaryType bx,by,bz; /* type of ghost nodes */
    DMDAStencilType st;
    DM            da;
} DMDALocalInfo;
```

Set values by stencil

```

DMDALocalInfo info;
 ierr = DMDAGetLocalInfo(grid,&info);CHKERRQ(ierr);
for (int j=info.ys; j<info.ys+info.ym; j++) {
  for (int i=info.xs; i<info.xs+info.xm; i++) {
    MatStencil row = {0}, col[5 ] = {{0}};
    PetscScalar v[ 5 ];
    PetscInt   ncols = 0;
    row.j      = j; row.i = i;
    // diagonal element:
    col[ncols].j = j; col[ncols].i = i; v[ncols++] = 4.;
    /* boundaries: top row */
    if (i>0)       {col[ncols].j = j;  col[ncols].i = i-1; v[ncols++] =
     ↪-1;};
    /* boundary left and right */
    if (j>0)       {col[ncols].j = j-1; col[ncols].i = i;  v[ncols++] =
     ↪-1;};
    if (j<info.my-1) {col[ncols].j = j+1; col[ncols].i = i;  v[ncols++] =
     ↪-1;};
  [et cetera]
```

DM Plex

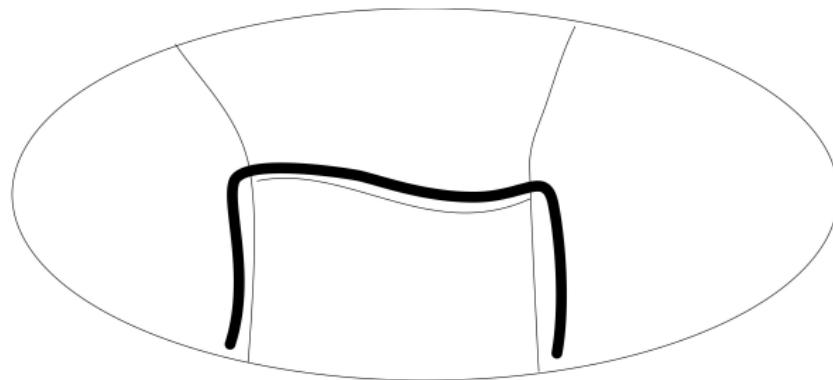
Support for unstructured grids and node/edge/cell relations.

This is complicated and under-documented.

IS and VecScatter: irregular grids

Irregular data movement

Example: collect distributed boundary onto a single processor
(this happens in the matrix-vector product)



Problem: figuring out communication is hard, actual communication is cheap

VecScatter

Preprocessing: determine mapping between input vector and output:

```
VecScatterCreate(Vec,IS,Vec,IS,VecScatter*)
// also Destroy
```

Application to specific vectors:

```
VecScatterBegin(VecScatter,
    Vec,Vec, InsertMode mode, ScatterMode direction)
VecScatterEnd (VecScatter,
    Vec,Vec, InsertMode mode, ScatterMode direction)
```

IS: index set

Index Set is a set of indices

```
ISCreateGeneral(comm,n,indices,PETSC_COPY_VALUES,&is);
/* indices can now be freed */
ISCreateStride (comm,n,first,step,&is);
ISCreateBlock (comm,bs,n,indices,&is);

ISDestroy(is);
```

Use MPI_COMM_SELF most of the time.

Various manipulations: ISSum, ISDifference,
ISInvertPermutations et cetera.

Also ISGetIndices / ISRestoreIndices / ISGetSize

Example: split odd and even

Input:

Process [0]

0.
1.
2.
3.
4.
5.

Process [1]

6.
7.
8.
9.
10.
11.

Output:

Process [0]

0.
2.
4.
6.
8.
10.

Process [1]

1.
3.
5.
7.
9.
11.

index sets for this example

```
// oddeven.c
IS oddeven;
if (procid==0) {
    ierr = ISCreateStride(comm,Nglobal/2,0,2,&oddeven); CHKERRQ(
        ↪ierr);
} else {
    ierr = ISCreateStride(comm,Nglobal/2,1,2,&oddeven); CHKERRQ(
        ↪ierr);
}
```

scatter for this example

```
VecScatter separate;
ierr = VecScatterCreate
  (in,oddeven,out,NULL,&separate); CHKERRQ(ierr);
ierr = VecScatterBegin
  (separate,in,out,INSERT_VALUES,SCATTER_FORWARD);
  ↪CHKERRQ(ierr);
ierr = VecScatterEnd
  (separate,in,out,INSERT_VALUES,SCATTER_FORWARD);
  ↪CHKERRQ(ierr);
```

Exercise 9 (oddeven)

Now alter the **IS** objects so that the output becomes:

Process [0]

10.
8.
6.
4.
2.
0.

Process [1]

11.
9.
7.
5.
3.
1.

Example: simulate allgather

```
/* create the distributed vector with one element per processor */
ierr = VecCreate(MPI_COMM_WORLD,&global);
ierr = VecSetType(global,VECMPI);
ierr = VecSetSizes(global,1,PETSC_DECIDE);

/* create the local copy */
ierr = VecCreate(MPI_COMM_SELF,&local);
ierr = VecSetType(local,VECSEQ);
ierr = VecSetSizes(local,ntids,ntids);
```

```
IS indices;  
ierr = ISCreateStride(MPI_COMM_SELF,ntids,0,1, &indices);  
/* create a scatter from the source indices to target */  
ierr = VecScatterCreate  
    (global,indices,local,indices,&scatter);  
/* index set is no longer needed */  
ierr = ISDestroy(&indices);
```

Example: even and odd indices

```
// oddeven.c
IS oddeven;
if (procid==0) {
    ierr = ISCreateStride(comm,Nglobal/2,0,2,&oddeven); CHKERRQ(
        ↪ierr);
} else {
    ierr = ISCreateStride(comm,Nglobal/2,1,2,&oddeven); CHKERRQ(
        ↪ierr);
}
```

scattering odd and even

```
VecScatter separate;
ierr = VecScatterCreate
  (in,oddeven,out,NULL,&separate); CHKERRQ(ierr);
ierr = VecScatterBegin
  (separate,in,out,INSERT_VALUES,SCATTER_FORWARD);
  ↪CHKERRQ(ierr);
ierr = VecScatterEnd
  (separate,in,out,INSERT_VALUES,SCATTER_FORWARD);
  ↪CHKERRQ(ierr);
```

SNES: Nonlinear solvers

Nonlinear problems

Basic equation

$$f(u) = 0$$

where u can be big, for instance nonlinear PDE.

Typical solution method:

$$u_{n+1} = u_n - J(u_n)^{-1} f(u_n)$$

Newton iteration.

Needed: function and Jacobian.

Basic SNES usage

User supplies function and Jacobian:

```
SNES          snes;
```

```
SNESCreate(PETSC_COMM_WORLD,&snes)
SNESSetType(snes,type);
SNESSetFromOptions(snes);
SNESDestroy(SNES snes);
```

where type:

- SNESLS Newton with line search
- SNESTR Newton with trust region
- several specialized ones

SNES specification: function evaluation

```
|| PetscErrorCode (*FunctionEvaluation)(SNES,Vec,Vec,void*);  
|| VecCreate(PETSC_COMM_WORLD,&r);  
|| SNESSetFunction(snes,r,FunctionEvaluation,*ctx);
```

SNES specification: jacobian evaluation

```
|| PetscErrorCode (*FormJacobian)(SNES,Vec,Mat,Mat,void*);  
|| MatCreate(PETSC_COMM_WORLD,&J);  
|| SNESSetJacobian(snes,J,J,FormJacobian,*ctx);
```

SNES solution

```
|| SNESolve(snes,/* rhs= */ PETSC_NULL,x)
|| SNESGetConvergedReason(snes,&reason)
|| SNESGetIterationNumber(snes,&its)
```

Example: two-variable problem

Define a context

```
typedef struct {
    Vec xloc,rloc; VecScatter scatter; } AppCtx;

/* User context */
AppCtx        user;

/* Work vectors in the user context */
VecCreateSeq(PETSC_COMM_SELF,2,&user.xloc);
VecDuplicate(user.xloc,&user.rloc);

/* Create the scatter between the global and local x */
ISCreateStride(MPI_COMM_SELF,2,0,1,&idx);
VecScatterCreate(x,idx,user.xloc,idx,&user.scatter);
```

In the user function:

```
PetscErrorCode FormFunction
  (SNES snes,Vec x,Vec f,void *ctx)
{
  VecScatterBegin(user->scatter,
  x,user->xloc,INSERT_VALUES,SCATTER_FORWARD); // & End
  VecGetArray(xloc,&xx);CHKERRQ(ierr);
  VecSetValue
    (f,0,/* something with xx[0]) & xx[1] */,
     INSERT_VALUES);
  VecRestoreArray(x,&xx);
  PetscFunctionReturn(0);
}
```

Jacobian calculation through finite differences

Jacobian calculation is difficult. It can be approximated through finite differences:

$$J(u)v \approx \frac{f(u + hv) - f(u)}{h}$$

```
MatCreateSNESMF(snes,&J);
SNESSet Jacobian
(snes,J,J,MatMFFDCompute Jacobian,(void*)&user);
```

Further possibilities

`SNESSetTolerances(SNES snes,double atol,double rtol,double stol, int its,int fcts);`

convergence test and monitoring, specific options for line search and trust region

adaptive convergence: `-snes_ksp_ew_conv` (Eisenstat Walker)

Solve customization

```
SNESSetType(snes,SNESTR); /* newton with trust region */  
SNESGetKSP(snes,&ksp)  
KSPGetPC(ksp,&pc)  
PCSetType(pc,PCNONE)  
KSPSetTolerances(ksp,1.e-4,PETSC_DEFAULT,PETSC_DEFAULT,
```

TS: Time stepping

Profiling, debugging

Basic profiling

- `-log_summary` flop counts and timings of all PETSc events
- `-info` all sorts of information, in particular

```
%% mpiexec yourprogram -info | grep malloc  
[0] MatAssemblyEnd_SeqAIJ():
```

Number of mallocs during `MatSetValues()` is 0

- `-log_trace` start and end of all events: good for hanging code

Log summary: overall

| | Max | Max/Min | Avg | Total |
|------------------|-----------|---------|-----------|-----------|
| Time (sec): | 5.493e-01 | 1.00006 | 5.493e-01 | |
| Objects: | 2.900e+01 | 1.00000 | 2.900e+01 | |
| Flops: | 1.373e+07 | 1.00000 | 1.373e+07 | 2.746e+07 |
| Flops/sec: | 2.499e+07 | 1.00006 | 2.499e+07 | 4.998e+07 |
| Memory: | 1.936e+06 | 1.00000 | | 3.871e+06 |
| MPI Messages: | 1.040e+02 | 1.00000 | 1.040e+02 | 2.080e+02 |
| MPI Msg Lengths: | 4.772e+05 | 1.00000 | 4.588e+03 | 9.544e+05 |
| MPI Reductions: | 1.450e+02 | 1.00000 | | |

Log summary: details

| | Max | Ratio | Max | Ratio | Avg | len | %T | %F | %M | %L | %R | %T | %F | %M | %L | %R | Mflop/s | |
|------------------|-----|-------|------------|-------|----------|-----|---------|-------|----|----|----|-------|----|----|----|-------|---------|-----|
| MatMult | 100 | 1.0 | 3.4934e-02 | 1.0 | 1.28e+08 | 1.0 | 8.0e+02 | 6 | 32 | 96 | 17 | 0 | 6 | 32 | 96 | 17 | 0 | 255 |
| MatSolve | 101 | 1.0 | 2.9381e-02 | 1.0 | 1.53e+08 | 1.0 | 0.0e+00 | 5 | 33 | 0 | 0 | 0 | 5 | 33 | 0 | 0 | 0 | 305 |
| MatLUFactorNum | 1 | 1.0 | 2.0621e-03 | 1.0 | 2.18e+07 | 1.0 | 0.0e+00 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 43 |
| MatAssemblyBegin | 1 | 1.0 | 2.8350e-03 | 1.1 | 0.00e+00 | 0.0 | 1.3e+05 | 0 | 0 | 3 | 83 | 1 | 0 | 0 | 3 | 83 | 1 | 0 |
| MatAssemblyEnd | 1 | 1.0 | 8.8258e-03 | 1.0 | 0.00e+00 | 0.0 | 4.0e+02 | 2 | 0 | 1 | 0 | 3 | 2 | 0 | 1 | 0 | 3 | 0 |
| VecDot | 101 | 1.0 | 8.3244e-03 | 1.2 | 1.43e+08 | 1.2 | 0.0e+00 | 1 | 7 | 0 | 0 | 35 | 1 | 7 | 0 | 0 | 35 | 243 |
| KSPSetup | 2 | 1.0 | 1.9123e-02 | 1.0 | 0.00e+00 | 0.0 | 0.0e+00 | 3 | 0 | 0 | 0 | 2 | 3 | 0 | 0 | 0 | 2 | 0 |
| KSPSolve | 1 | 1.0 | 1.4158e-01 | 1.0 | 9.70e+07 | 1.0 | 8.0e+02 | 26100 | 96 | 17 | 92 | 26100 | 96 | 17 | 92 | 26100 | 194 | |

User events

```
#include "petsclog.h"
int USER EVENT;
PetscLogEventRegister(&USER EVENT,"User event name",0);
PetscLogEventBegin(USER EVENT,0,0,0,0);
/* application code segment to monitor */
PetscLogFlops(number of flops for this code segment);
PetscLogEventEnd(USER EVENT,0,0,0,0);
```

Program stages

```
|| PetscLogStagePush(int stage); /* 0 <= stage <= 9 */
|| PetscLogStagePop();
|| PetscLogStageRegister(int stage,char *name)
```

Debugging

- Use of CHKERRQ and SETERRQ for catching and generating error
- Use of PetscMalloc and PetscFree to catch memory problems;
CHKMEMQ for instantaneous memory test (debug mode only)
- Better than PetscMalloc: PetscMalloc1 aligned to PETSC_MEMALIGN